



Lesson #13

Logistic Regression

- Classification
- Binary Classification
- Decision Boundary
- Cost Function
- Multiclass Classification
- Regularization (L1, L2)
- Hands on Scikit-Learn



MEDICAL MODEL



HEALTHY



SICK



SPAM CLASSIFIER MODEL



NOT SPAM



SPAM

Classification Problem


3





Test



Grades

Student 1
Test: 9/10 
Grades: 8/10

Student 2
Test: 3/10 
Grades: 4/10


Student 3
Test: 7/10 
Grades: 6/10

Binary Classification Problem

- Email = {spam, not spam}
- Medical model = {healthy, sick}
- Fraudulent operation = {yes, not}
- Academic acceptance = {success, fail}
- Movie review = {good, bad}

$Y \in \{0,1\}$  0: negative class
1: positive class

Binary Classification Problem (observation #1)



What if??



Threshold classifier output as $h_{\theta}(x)$:

- If $h_{\theta}(x) \geq 0.5$, predict $y = 1$
- If $h_{\theta}(x) < 0.5$, predict $y = 0$

Binary Classification Problem (observation #2)

- Y assume only two values: 0 or 1.
- In linear case, $h_{\theta}(x) \geq 1$ and $h_{\theta}(x) \leq 0$ can occur.



Logistic Regression
 $0 \leq h_{\theta}(x) \leq 1$

Logistic Regression - Hypothesis Representation


Target $\rightarrow 0 \leq h_{\theta}(x) \leq 1$

$$h_{\theta}(x) = \theta^T x \quad (\text{doesn't work})$$

$$h_{\theta}(x) = g(z) \quad , \text{where } z = \theta^T x$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

Sigmoid function or
Logistic function



Logistic Regression - Decision Boundary

Suppose:

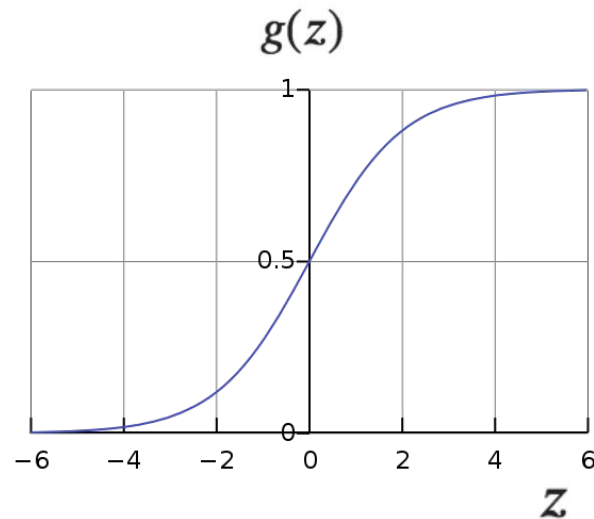
Predict $y = 1$ if $h_{\theta}(x) \geq 0.5$

$$g(z) \geq 0.5 \text{ when } z \geq 0$$

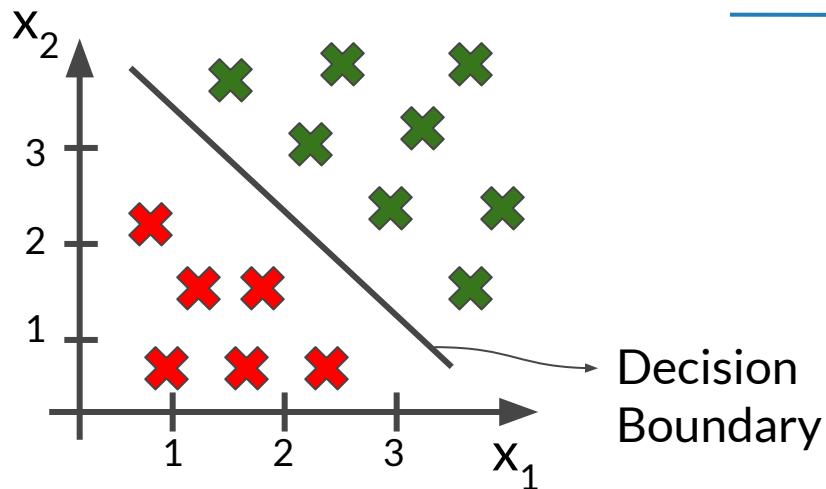
Suppose:

Predict $y = 0$ if $h_{\theta}(x) < 0.5$

$$g(z) < 0.5 \text{ when } z < 0$$



Logistic Regression - Decision Boundary



Suppose:

Predict $y = 1$ if $z \geq 0$

$$-3 + x_1 + x_2 \geq 0$$

$$x_1 + x_2 \geq 3$$

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

$$\theta = \begin{bmatrix} -3 \\ 1 \\ 1 \end{bmatrix}$$

$$z = -3 + x_1 + x_2$$

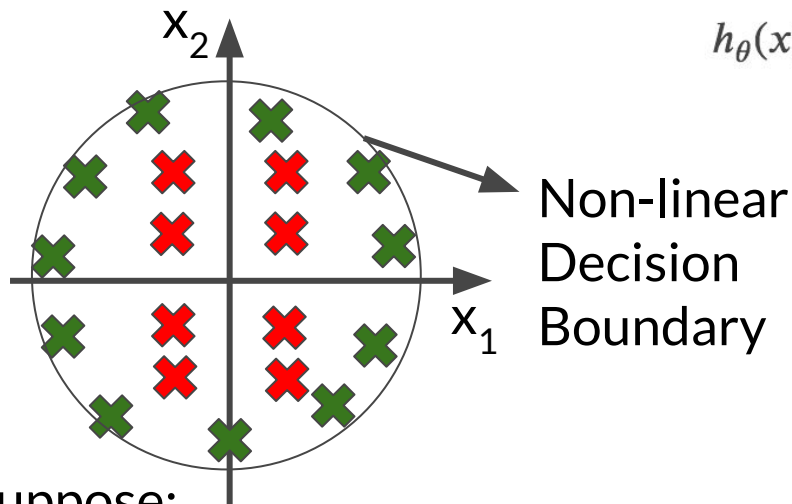
Suppose:

Predict $y = 0$ if $z < 0$

$$-3 + x_1 + x_2 < 0$$

$$x_1 + x_2 < 3$$

Logistic Regression - Decision Boundary



Suppose:

Predict $y = 1$ if $z \geq 0$

$$\begin{aligned} -1 + x_1^2 + x_2^2 &\geq 0 \\ x_1^2 + x_2^2 &\geq 1 \end{aligned}$$

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

$$\theta = \begin{bmatrix} -1 \\ 0 \\ 0 \\ 1 \\ 1 \end{bmatrix}$$

$$z = -1 + x_1^2 + x_2^2$$

Suppose:



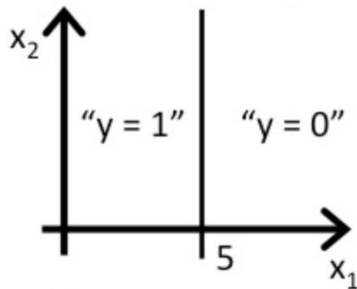
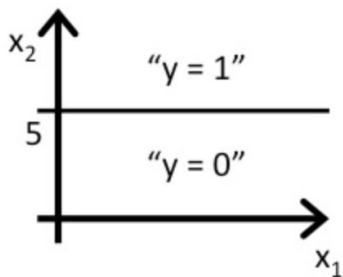
Predict $y = 0$ if $z < 0$

$$\begin{aligned} -1 + x_1^2 + x_2^2 &< 0 \\ x_1^2 + x_2^2 &< 1 \end{aligned}$$

Logistic Regression - Decision Boundary

Consider logistic regression with two features x_1 and x_2 . Suppose $\Theta_0 = 5$, $\Theta_1 = -1$ and $\Theta_2 = 0$, so that $h_{\Theta}(x) = g(5 - x_1)$.


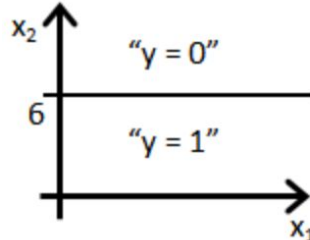
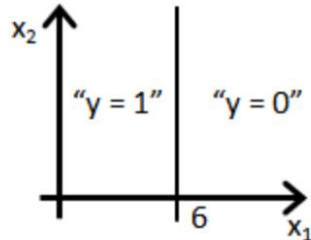
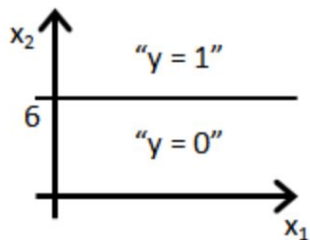
Which of these shows the decision boundary of $h_{\Theta}(x)$?



Logistic Regression - Decision Boundary

Consider logistic regression with two features x_1 and x_2 . Suppose $\Theta_0 = 6$, $\Theta_1 = 0$ and $\Theta_2 = -1$, so that $h_{\Theta}(x) = g(\Theta_0 + \Theta_1 x_1 + \Theta_2 x_2)$.

Which of these shows the decision boundary of $h_{\Theta}(x)$?



Logistic Regression - Decision Boundary

Suppose that you have trained a logistic regression classifier, and it outputs on a new example x a prediction $h_{\theta}(x) = 0.4$. This means (check all that apply):

Our estimate for $P(y = 0|x; \theta)$ is 0.4.

Our estimate for $P(y = 0|x; \theta)$ is 0.6.

Our estimate for $P(y = 1|x; \theta)$ is 0.4.

Our estimate for $P(y = 1|x; \theta)$ is 0.6.

RECAP

 $f(x)$

cost function

Training Set: $\{(\mathbf{x}^1, y^1), (\mathbf{x}^2, y^2), \dots, (\mathbf{x}^m, y^m)\}$
m examples

$$x \in \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix}, \text{ <n+1 elements> } , x_0 = 1, y \in \{0, 1\}$$

$$h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$$

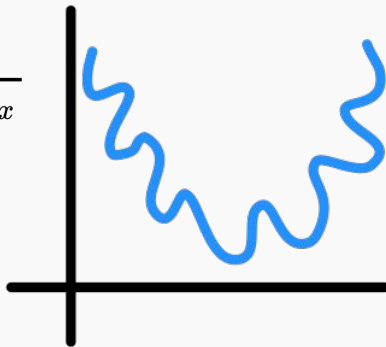
How to fit the parameter θ ?

Cost Function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^i) - y^i)^2 \rightarrow \text{cost}(h_{\theta}(x), y)$$

non-convex

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$




convex

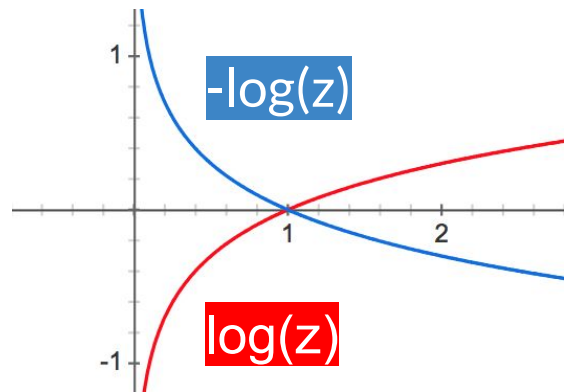


Logistic Regression Cost Function

$$\text{cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y=1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y=0 \end{cases}$$




$$\begin{array}{ll} h_{\theta}(x) \rightarrow 1 & h_{\theta}(x) \rightarrow 0 \\ \text{cost} \rightarrow 0 & \text{cost} \rightarrow \infty \end{array}$$



Logistic Regression Cost Function

$$\text{cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y=1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y=0 \end{cases}$$

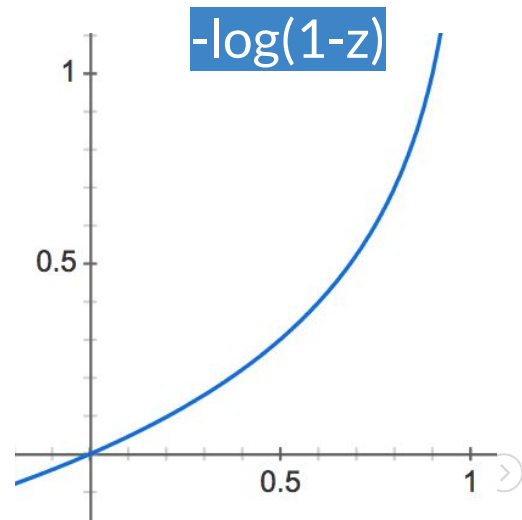


$$h_{\theta}(x) \rightarrow 1$$

$$\text{cost} \rightarrow \infty$$

$$h_{\theta}(x) \rightarrow 0$$

$$\text{cost} \rightarrow 0$$



Simplified Cost Function & Gradient Descent

Logistic Regression Cost Function

$$\text{cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y=1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y=0 \end{cases}$$

$$\text{cost}(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$$

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))]$$

Cost Function - Vectorized Implementation

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))]$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_k \end{bmatrix}$$

$$h = g(X\theta)$$

[m; k+1] x [k+1; 1] = [m; 1]

$$J(\theta) = \frac{1}{m} \cdot \left(-y^T \log(h) - (1 - y)^T \log(1 - h) \right)$$

[1; m] x [m; 1] = scalar



$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y)(\log(1 - h_{\theta}(x^{(i)})))]$$

```
def cost_function(theta, X, y):  
    thetaX = logistic(np.matmul(X, theta))  
    return -1/len(y) * np.sum(y*np.log(thetaX) + (1-y)*np.log(1 - thetaX))
```

General Form of Gradient Descent

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

Repeat {

$$\theta_j := \theta_j - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}



Vectorized Implementation

$$\theta := \theta - \frac{\alpha}{m} X^T (g(X\theta) - \vec{y})$$


gradient

$$\theta := \theta - \frac{\alpha}{m} X^T (g(X\theta) - \vec{y})$$

```
def gradient_descent_multi(theta_, X, y, alpha, iterations):  
    m = len(X)  
    theta = theta_.copy()  
    cost_history = []  
    for i in range(iterations):  
        gradient = (1/m) * np.dot(X.T, logistic(np.dot(X, theta)) - y)  
        theta = theta - (alpha * gradient)  
        cost_history.append(cost_function(theta, X, y))  
    return theta, cost_history
```

Admitted vs not Admitted

25



```
# define X and y
X = np.column_stack((np.ones(data.shape[0]),
                      data[["X1_scaled", "X2_scaled"]]))
y = data.Admitted.astype(np.int64).values.reshape(-1,1)

# define m and n
m,n = X.shape


# guess an initial value for theta
theta = np.zeros((n,1))
```



```
# using gradient descent  
# theta, X, y, alpha, iterations  
theta_batch, cost_history = gradient_descent_multi(theta,X,y,1,400)
```

```
array([[1.65947664],  
       [3.8670477 ],  
       [3.60347302]])
```


Cost Function vs #iterations (using gradient descent)



Lesson#13 - Some experimentation.ipynb

Up to section 2.10

- Plotting the decision boundary
- Learning parameters using `fmin_tnc`



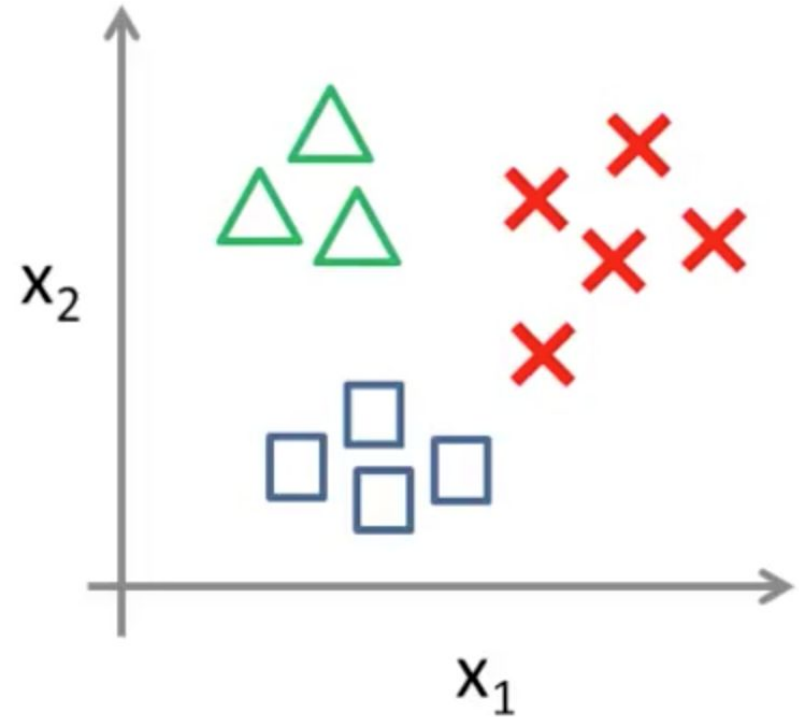
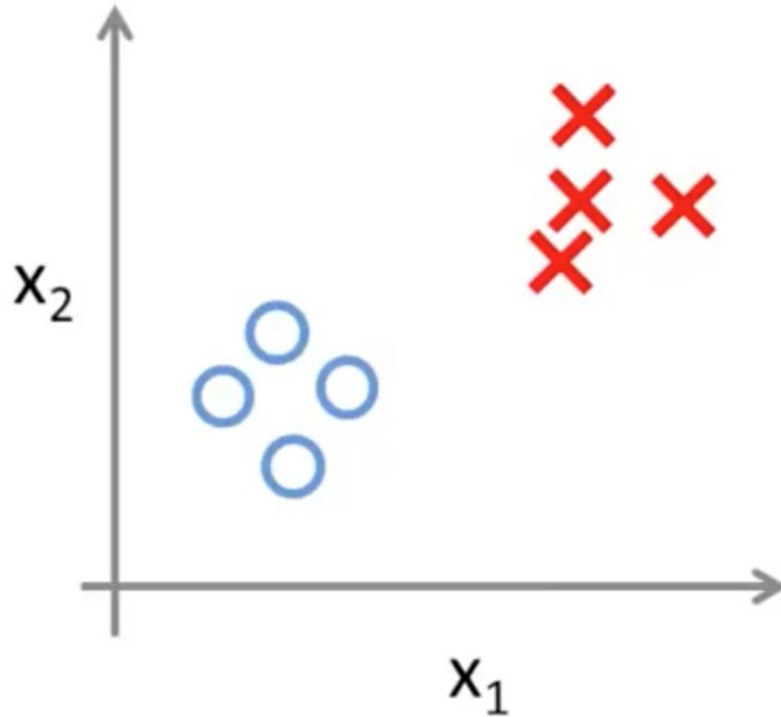
Multiclass Classification:

One vs All

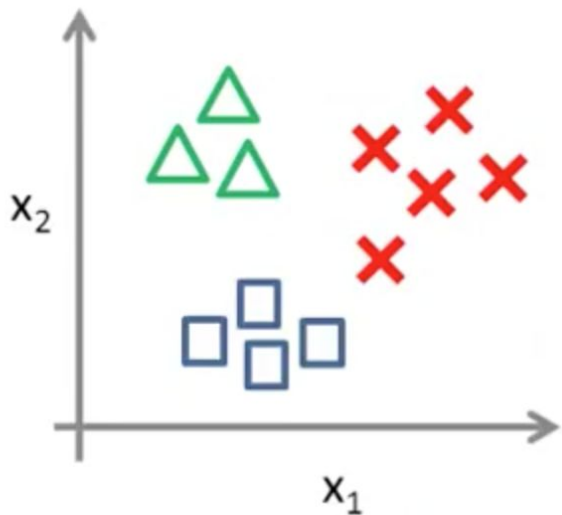
Multiclass Classification


- Email foldering/tagging: work, ad, family, friends, hobby
- Medical diagrams: not ill, cold, flu
- Weather: sunny, cloudy, rain, snow

Binary vs Multiclass Classification






Multiclass Classification (One vs All)



Class 1: 

Class 2: 

Class 3: 



Multiclass Classification (One vs All)

Train a logistic regression classifier $h_{\theta}^{(i)}(x)$ for each class i :


$$h_{\theta}^i(x) = P(y = i|x; \theta)$$

On a new input x , to make a prediction, pick the class i that maximizes:

$$\max_i h_{\theta}^{(i)}(x)$$

overfitting problem


Logistic Regression




$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

(g = sigmoid function)

Underfit



$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2 + \theta_5 x_1 x_2)$$



$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_1^2 + \theta_3 x_1^2 x_2 + \theta_4 x_1^2 x_2^2 + \theta_5 x_1^2 x_2^3 + \theta_6 x_1^3 x_2 + \dots)$$

Overfit

```
import numpy as np
from sklearn.preprocessing import PolynomialFeatures

a = np.arange(6).reshape(3, 2)
print(a)
[[0 1]
 [2 3]
 [4 5]]

poly = PolynomialFeatures(2)
poly.fit_transform(a)

array([[ 1.,  0.,  1.,  0.,  0.,  1.],
       [ 1.,  2.,  3.,  4.,  6.,  9.],
       [ 1.,  4.,  5., 16., 20., 25.]])
```

$$terms = \binom{n+d}{d} = \binom{2+2}{2} = 6$$

$$[1, x_1, x_2, x_1^2, x_1 x_2, x_2^2]$$

$$\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_1 x_2 + \theta_5 x_2^2$$


Addressing Overfitting

1. Reduce number of features
 - a. Manually select which feature to keep
2. Regularization
 - a. Keep all the features, but reduce magnitude/values of parameters Θ_j
 - b. Works well when we have a lot of features, each of which contributes a bit to predicting y .


Intuition - Regularized Linear Regression

$$\min_{\theta} \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

Regularization
Parameter



$$\theta_0 + \theta_1 x + \theta_2 x^2$$



$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$

Suppose we penalize and make Θ_3 and Θ_4 very small

$$\min_{\theta} \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + 1000 \cdot \theta_3^2 + 1000 \cdot \theta_4^2$$


Intuition - Gradient Descent (L2 - Ridge)

Repeat {


$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \left[\left(\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{m} \theta_j \right] \quad j \in \{1, 2, \dots, n\}$$

}


$$\theta_j := \theta_j (1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Intuition - Regularized Logistic Regression




$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_1^2 + \theta_3 x_1^2 x_2 + \theta_4 x_1^2 x_2^2 + \theta_5 x_1^2 x_2^3 + \dots)$$

Cost function:

$$J(\theta) = - \left[\frac{1}{m} \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$

Intuition - Regularized Logistic Regression (L2)



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_1^2 + \theta_3 x_1^2 x_2 + \theta_4 x_1^2 x_2^2 + \theta_5 x_1^2 x_2^3 + \dots)$$

Cost function:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

L2 - Ridge Regression

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y)(\log(1 - h_{\theta}(x^{(i)})))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

$$\theta_i : i = 0 \dots n$$

```
def cost_function_reg_l2(theta, X, y, lambda_):  
    m = len(y)  
    thetaX = logistic(np.matmul(X, theta))  
    regularization = lambda_/(2*m) * np.sum(theta[1:]**2)  
    return -1/m * np.sum(y*np.log(thetaX) + (1-y)*np.log(1 - thetaX)) + regularization
```

L1 - Lasso Regression

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y)(\log(1 - h_{\theta}(x^{(i)})))] + \frac{\lambda}{2m} \sum_{j=1}^n |\theta_j|$$

$$\theta_i : i = 0 \dots n$$

```
def cost_function_reg_l1(theta, X, y, lambda_):  
    m = len(y)  
    thetaX = logistic(np.matmul(X, theta))  
    regularization = lambda_/(2*m) * np.sum(np.absolute(theta[1:]))  
    return -1/m * np.sum(y*np.log(thetaX) + (1-y)*np.log(1 - thetaX)) + regularization
```


Intuition - Gradient Descent (L2)





Repeat {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \left[\left(\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{m} \theta_j \right] \quad j \in \{1, 2 \dots n\}$$

}


$$\theta_j := \theta_j (1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$



Lesson#13 - Some experimentation.ipynb




- Visualize different decision boundaries
- Analyze the results (thetas) considering the follow regularization techniques: L1, L2
- Changes values of lambda, iterations, so on
- Check the equations




Binary Classification

admit	gpa	gre
0	3.177277	594.102992
0	3.412655	631.528607
0	2.728097	553.714399
0	3.093559	551.089985
0	3.141923	537.184894

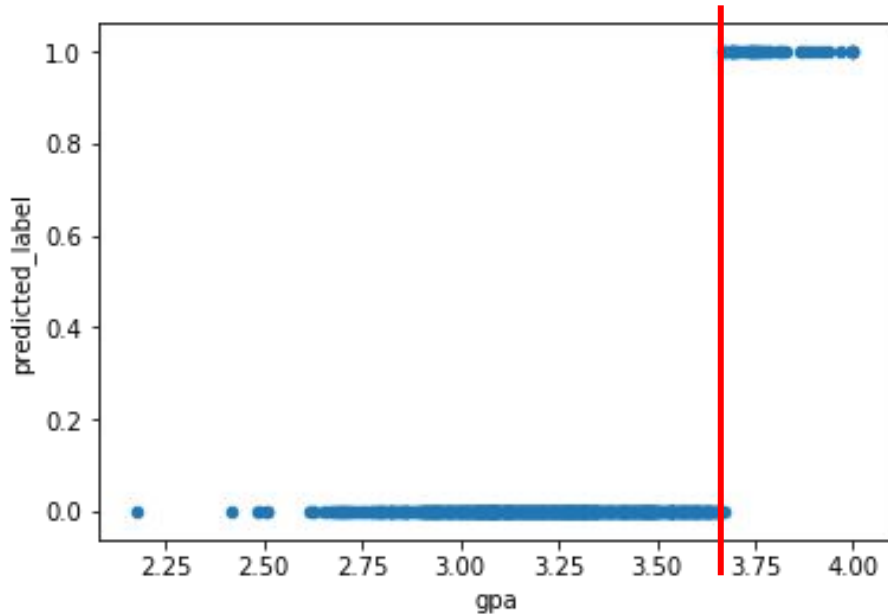


Logistic Regression Model (fit, predict prob.)


```
# create a model
model = LogisticRegression()
# training
model.fit(data[["gpa"]], data["admit"])
# predict
prob = model.predict_proba(data[["gpa"]])
# store in a dataframe
data["P(Admitted)"] = prob[:,1]
```



Logistic Regression Model (fit, predict class)



```
model.predict(data[["gpa"]])
```



```
model.predict_proba(data[["gpa"]])[:,1]
```

Evaluating Binary Classifiers

Prediction	Observation	
	Admitted (1)	Rejected (0)
Admitted (1)	True Positive (TP)	False Positive (FP)
Rejected (0)	False Negative (FN)	True Negative (TN)

```
from sklearn.metrics import confusion_matrix
```

```
tn, fp, fn, tp = confusion_matrix(data.admit,  
                                  data.predicted_label).ravel()
```

Evaluating Binary Classifiers

Prediction	Observation	
	Admitted (1)	Rejected (0)
Admitted (1)	True Positive (TP)	False Positive (FP)
Rejected (0)	False Negative (FN)	True Negative (TN)

$$Accuracy = \frac{\text{\#correct predictions}}{\text{\#observations}}$$

$$Accuracy = \frac{TP + TN}{TP + TN + FN + FP}$$

$$TPR = \frac{\text{\#true positives}}{\text{\#true positives} + \text{\#false negatives}}$$

Recall

$$TNR = \frac{\text{\#true negatives}}{\text{\#true negatives} + \text{\#false positives}}$$

Evaluating Binary Classifiers

Prediction	Observation	
	Admitted (1)	Rejected (0)
Admitted (1)	True Positive (TP)	False Positive (FP)
Rejected (0)	False Negative (FN)	True Negative (TN)

$$precision = \frac{TP}{(TP + FP)}$$

$$precision = \frac{TN}{(TN + FN)}$$

$$F_1 = 2 \times \frac{Precision \times Recall}{Precision + Recall}$$

```
from sklearn.metrics import classification_report

print(classification_report(data[["admit"]],
                             data[["predicted_label"]],
                             target_names=['0', '1']))
```

	precision	recall	f1-score	support
0	0.64	0.96	0.77	400
1	0.67	0.13	0.21	244
accuracy			0.65	644
macro avg	0.66	0.54	0.49	644
weighted avg	0.66	0.65	0.56	644

Multiclass Classification

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin
0	18.0	8	307.0	130.0	3504.0	12.0	70	1
1	15.0	8	350.0	165.0	3693.0	11.5	70	1
2	18.0	8	318.0	150.0	3436.0	11.0	70	1
3	16.0	8	304.0	150.0	3433.0	12.0	70	1
4	17.0	8	302.0	140.0	3449.0	10.5	70	1

origin -- Integer and Categorical. 1: North America, 2: Europe, 3: Asia.

Dummy Variables

cyl_3	cyl_4	cyl_5	cyl_6	cyl_8
0	0	0	0	1
0	0	0	0	1
0	0	0	0	1
0	0	0	0	1
0	0	0	0	1

```
dummy_cylinders = pd.get_dummies(
    cars["cylinders"], prefix="cyl")
cars = pd.concat([cars, dummy_cylinders],
                  axis=1)
cars.head()
```

[illegible]

Training a Multiclass Logistic Regression Model

```
from sklearn.linear_model import LogisticRegression

unique_origins = cars["origin"].unique()
unique_origins.sort()

models = {}
features = [c for c in train.columns
             if c.startswith("cyl") or c.startswith("year")]

for origin in unique_origins:
    model = LogisticRegression()

    X_train = train[features]
    y_train = train["origin"] == origin

    model.fit(X_train, y_train)
    models[origin] = model
```

Testing (One vs All)

	1	2	3
0	0.613723	0.131164	0.262305
1	0.536781	0.226177	0.236130
2	0.613723	0.131164	0.262305
3	0.678392	0.174871	0.154612
4	0.616443	0.226177	0.162931

